

# E-Graphs and Equality Saturation (in Haskell)

Rodrigo Mesquita

March 8, 2023

# E-graphs and eq-sat are cool

Published applications of equality saturation and e-graphs

- Automatic vectorization of digital signal processing code
- Tensor graph superoptimization
- Algebraic metaprogramming and symbolic computation
- and more...

# E-graphs and eq-sat are cool

And perhaps in the near future...

- A symbolic mathematics library in Haskell
- Pattern-match coverage checking in GHC
- <Insert your idea here>

# What are e-graphs?

## Definition (E-graphs)

- An e-graph is a data structure that
  - compactly represents *equivalence classes* of expressions
  - while maintaining a key invariant: the equivalence relation is closed under *congruence*<sup>a</sup>
- Concretely, an e-graph is a set of equivalence classes.
  - Each e-class is a set of equivalent e-nodes
  - An e-node represents a expression from a given language (e.g.  $x * 1$ )
  - An e-node is a function symbol paired with a list of children e-classes  $f(c_1, c_2, \dots)$

---

<sup>a</sup>Intuitively,  $a \equiv b \Rightarrow f(a) \equiv f(b)$

# What are e-graphs?

## Definition (E-graphs)

- An e-graph is a data structure that
  - compactly represents *equivalence classes* of expressions
  - while maintaining a key invariant: the equivalence relation is closed under *congruence*<sup>a</sup>
- Concretely, an e-graph is a set of equivalence classes.
  - Each e-class is a set of equivalent e-nodes
  - An e-node represents a expression from a given language (e.g.  $x * 1$ )
  - An e-node is a function symbol paired with a list of children e-classes  $f(c_1, c_2, \dots)$

---

<sup>a</sup>Intuitively,  $a \equiv b \Rightarrow f(a) \equiv f(b)$

# What are e-graphs?

## Definition (E-graphs)

- An e-graph is a data structure that
  - compactly represents *equivalence classes* of expressions
  - while maintaining a key invariant: the equivalence relation is closed under *congruence*<sup>a</sup>
- Concretely, an e-graph is a set of equivalence classes.
  - Each e-class is a set of equivalent e-nodes
  - An e-node represents a expression from a given language (e.g.  $x * 1$ )
  - An e-node is a function symbol paired with a list of children e-classes  $f(c_1, c_2, \dots)$

---

<sup>a</sup>Intuitively,  $a \equiv b \Rightarrow f(a) \equiv f(b)$

# What are e-graphs?

## Definition (E-graphs)

- An e-graph is a data structure that
  - compactly represents *equivalence classes* of expressions
  - while maintaining a key invariant: the equivalence relation is closed under *congruence*<sup>a</sup>
- Concretely, an e-graph is a set of equivalence classes.
  - Each e-class is a set of equivalent e-nodes
  - An e-node represents a expression from a given language (e.g.  $x * 1$ )
  - An e-node is a function symbol paired with a list of children e-classes  $f(c_1, c_2, \dots)$

---

<sup>a</sup>Intuitively,  $a \equiv b \Rightarrow f(a) \equiv f(b)$

# What are e-graphs?

## Definition (E-graphs)

- An e-graph is a data structure that
  - compactly represents *equivalence classes* of expressions
  - while maintaining a key invariant: the equivalence relation is closed under *congruence*<sup>a</sup>
- Concretely, an e-graph is a set of equivalence classes.
  - Each e-class is a set of equivalent e-nodes
  - An e-node represents a expression from a given language (e.g.  $x * 1$ )
  - An e-node is a function symbol paired with a list of children e-classes  $f(c_1, c_2, \dots)$

---


<sup>a</sup>Intuitively,  $a \equiv b \Rightarrow f(a) \equiv f(b)$



# E-graphs represent congruence relations over expressions

## Example (E-graph)


- This e-graph represents  $(a * 2)/2$
- In fact, it represents  $a$ ,  $2$  and  $a * 2$  too



# E-graphs represent congruence relations over expressions

## Example (E-graph)


- This e-graph represents  $(a * 2)/2$
- In fact, it represents  $a$ ,  $2$  and  $a * 2$  too



# E-graphs represent congruence relations over expressions

## Example (E-graph)


- Now we **represent** the expression  $a \ll 1$  in a new e-class
- And **merge** that e-class with the e-class representing  $a * 2$ 
  - We merge these two classes because they represent equivalent expressions
  - By congruence,  $(a * 2)/2$  is now seen as equivalent to  $(a \ll 1)/2$



# E-graphs represent congruence relations over expressions

## Example (E-graph)


- Now we **represent** the expression  $a \ll 1$  in a new e-class
- And **merge** that e-class with the e-class representing  $a * 2$ 
  - We merge these two classes because they represent equivalent expressions
  - By congruence,  $(a * 2)/2$  is now seen as equivalent to  $(a \ll 1)/2$



# E-graphs represent congruence relations over expressions

## Example (E-graph)

- Now we **represent** the expression  $a \ll 1$  in a new e-class
- And **merge** that e-class with the e-class representing  $a * 2$ 
  - We merge these two classes because they represent equivalent expressions
  - By congruence,  $(a * 2)/2$  is now seen as equivalent to  $(a \ll 1)/2$



# What about equality saturation?

## Definition (Equality Saturation)

- Eq-sat is a technique that leverages e-graphs to implement term rewriting systems
- In short, we
  - Represent an expression in the e-graph
  - Repeatedly apply pattern-based rewrites until saturation (rewrite rules like  $x * 1 \rightarrow x$  where  $x$  matches any expression)
  - Extract the best expression equivalent to the input

# What about equality saturation?

## Definition (Equality Saturation)

- Eq-sat is a technique that leverages e-graphs to implement term rewriting systems
- In short, we
  - Represent an expression in the e-graph
  - Repeatedly apply pattern-based rewrites until saturation (rewrite rules like  $x * 1 \rightarrow x$  where  $x$  matches any expression)
  - Extract the best expression equivalent to the input

# What about equality saturation?

## Definition (Equality Saturation)

- Eq-sat is a technique that leverages e-graphs to implement term rewriting systems
- In short, we
  - Represent an expression in the e-graph
  - Repeatedly apply pattern-based rewrites until saturation (rewrite rules like  $x * 1 \rightarrow x$  where  $x$  matches any expression)
  - Extract the best expression equivalent to the input



# What about equality saturation?


## Definition (Equality Saturation)

- Eq-sat is a technique that leverages e-graphs to implement term rewriting systems
- In short, we
  - Represent an expression in the e-graph
  - Repeatedly apply pattern-based rewrites until saturation (rewrite rules like  $x * 1 \rightarrow x$  where  $x$  matches any expression)
  - Extract the best expression equivalent to the input

## Example (Rewriting $(a * 2)/2$ )

Rewrite rules:

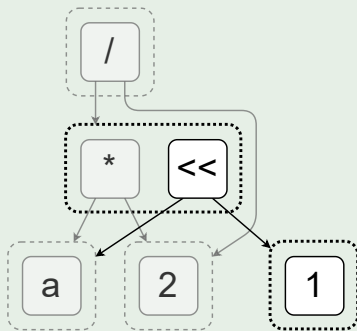
- $a * 2 = a \ll 1$
- $(a * b)/c = a * (b/c)$
- $x/x = 1$
- $x * 1 = x$



## Example (Rewriting $(a * 2)/2$ )

Rewrite rules:


- $a * 2 = a \ll 1$
- $(a * b)/c = a * (b/c)$
- $x/x = 1$
- $x * 1 = x$



## Example (Rewriting $(a * 2)/2$ )

Rewrite rules:

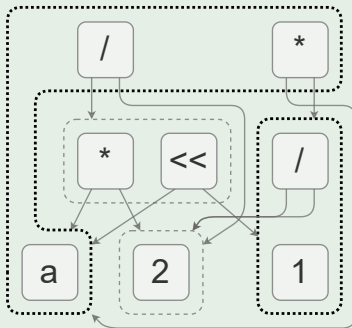
- $a * 2 = a \ll 1$
- $(a * b)/c = a * (b/c)$
- $x/x = 1$
- $x * 1 = x$



## Example (Rewriting $(a * 2)/2$ )

Rewrite rules:

- $a * 2 = a \ll 1$
- $(a * b)/c = a * (b/c)$
- $x/x = 1$
- $x * 1 = x$



# Fast and extensible equality saturation<sup>1</sup>

- Invariant restoration technique called *rebuilding*
- Mechanism called *e-class analysis* to integrate domain specific analysis into the e-graph

---

<sup>1</sup>egg: Fast and Extensible Equality Saturation (POPL 2021)

# Rebuilding

The key idea: defer e-graph invariant restoration to the call of *rebuild*

- Under certain workloads (automated theorem provers such as Z3), e-graph invariants are restored after every operation
- For equality saturation, we can greatly benefit from deferred invariant maintenance:
  - Search for all pattern matches
  - Apply all rewrites that matched<sup>2</sup>
  - Restore invariants<sup>3</sup>

---

<sup>2</sup>We can't search patterns while applying rewrites – the invariants aren't maintained w/out **rebuild**

<sup>3</sup>Potentially saves work because we deduplicate the *worklist*

# Rebuilding

The key idea: defer e-graph invariant restoration to the call of *rebuild*

- Under certain workloads (automated theorem provers such as Z3), e-graph invariants are restored after every operation
- For equality saturation, we can greatly benefit from deferred invariant maintenance:
  - Search for all pattern matches
  - Apply all rewrites that matched<sup>2</sup>
  - Restore invariants<sup>3</sup>

---

<sup>2</sup>We can't search patterns while applying rewrites – the invariants aren't maintained w/out **rebuild**

<sup>3</sup>Potentially saves work because we deduplicate the *worklist*



# E-class Analysis

The key idea: attach data to each e-class and use it to modify the e-graph

- Purely syntactic rewrites are insufficient in certain applications, in which we need domain knowledge (e.g. constant folding)
- Previously, this knowledge was added by modifying the e-graph with ad-hoc passes
- E-class analysis allows the expression of analysis over the e-graph
  - Each e-class has data from a semilattice
  - Merging e-classes merges the data
  - We can use the data to modify the e-graph
  - (See interface of Analysis)

# E-class Analysis

The key idea: attach data to each e-class and use it to modify the e-graph

- Purely syntactic rewrites are insufficient in certain applications, in which we need domain knowledge (e.g. constant folding)
- Previously, this knowledge was added by modifying the e-graph with ad-hoc passes
- E-class analysis allows the expression of analysis over the e-graph
  - Each e-class has data from a semilattice
  - Merging e-classes merges the data
  - We can use the data to modify the e-graph
  - (See interface of Analysis)

# E-class Analysis

The key idea: attach data to each e-class and use it to modify the e-graph

- Purely syntactic rewrites are insufficient in certain applications, in which we need domain knowledge (e.g. constant folding)
- Previously, this knowledge was added by modifying the e-graph with ad-hoc passes
- E-class analysis allows the expression of analysis over the e-graph
  - Each e-class has data from a semilattice
  - Merging e-classes merges the data
  - We can use the data to modify the e-graph
  - (See interface of Analysis)

# E-graphs in hegg (Demo)

## Remark

Follow through with the documentation: `Data.Equality.Graph`

- Why are expressions in their functorial form?

```
data SymExpr a
  = Const Double
  | Symbol String
  | a : + : a
  | a : * : a
  | a : / : a
deriving (Functor, Foldable, Traversable)
```

# References

- egg: Fast and Extensible Equality Saturation
- Relational E-matching